

CSE443/543: High Performance Computing

Exercise #19: Collective communication

Points: 40

Objective: The objective of this exercise is to:

- Gain familiarity with collective communication.
- Understand the use of MPI collective communication operations.
- Learn to run MPI programs via SLURM on a compute cluster.

Submission: Upload the following at the end of the lab exercise via **Canvas CODE plugin**:

1. This MS-Word document saved as a PDF file with the convention MUID_Exercise19.pdf.
2. The program you completed as part of this exercise with the source file named with the convention MUID_exercise19.cpp.

Part #0: Review online lecture



Required video review:

Prior to working on this exercise, ensure you review the lecture video on collective communications on Canvas.

Part #1: Short answer questions

Provide a brief (2-to-3 sentences) response to each of the following questions.

1. What is a virtual synchronization point? Explain with a suitable MPI call. How is it different from a barrier?
 - a. 2 Advantages

Barrier is waiting for all the processes to finish. Virtual synchronization point – all the processes happen to use the same program, they do not really synchronize but just virtually synchronize.

2. Briefly describe 2 significant differences between conceptual broadcast versus scatter operations

Broadcast	Scatter
One sender which sends exactly the same data to all the processes	One sender process that send different data to all the processes
Sender does not send data to itself	Sender sends data to itself

- Given the following MPI code fragment from process with rank 0, complete the complementary collective operation on other processes

```
void doManagerTasks(const std::string& data) {  
    int strSize = data.size() + 1;  
    MPI_Bcast(&strSize, 1, MPI_INT, 0, MPI_COMM_WORLD);  
    MPI_Bcast(&data[0], size, MPI_CHAR, 0, MPI_COMM_WORLD);  
}
```

```
std::string recvData() {  
    mpi::broadcast(&data[0], size, MPI_CHAR, 0, MPI_COMM_WORLD);  
}
```


Part #2: Programming with collective communication

In this part of the exercise, you will be required to complete 3 interview/exam style questions in the supplied `exercise19.cpp` starter code.

Part #2.1: Setting up VS-Code project

Estimated time to complete: 5 minutes

- Log into OSC's OnDemand portal via <https://ondemand.osc.edu/>. Login with your OSC id and password.
- Startup a VS-Code server and connect to VS-Code. Ensure you switch to your workspace. Your VS-Code window should appear as shown in the adjacent screenshot.
- Next, create a new VS-Code project in the following manner:
 - Start a new terminal in VS-Code
 - In the VS-Code terminal use the following commands:



```
$ # First change to your workspace directory  
$ cd ~/cse443  
$ # Use ls to check if workspace.code-workspace file is in pwd  
$ # Next copy the basic template for a C++ project  
$ cp -r /fs/ess/PMIU0184/cse443/templates/mpi exercise19  
$ # Copy the starter code for this exercise  
$ cp /fs/ess/PMIU0184/cse443/exercises/exercise19/* exercise19
```

Part #2.2: Mean and variance

In the supplied `exercise19.cpp` starter code, implement the `getMeanAndVar` method. The `getMeanAndVar` method is called on `n` processes of an MPI program. Each process will have a different value (as shown in the figure below). This method must be implemented to compute and return the mean and variance on every process (as shown in the figure below). **Note that your implementation must use only collective communication and computation operations.** Hint: `all_reduce`.

mean (average): μ

$$\mu = \frac{\sum_{i=1}^n x_i}{n}$$

variance: σ^2

$$\sigma^2 = \frac{\sum_{i=1}^n (x_i - \mu)^2}{n}$$

Rank:	0	1	2
value (the parameter to getMeanAndVar)	1	2	3
Return values	{2.5, 0.66}	{2.5, 0.66}	{2.5, 0.66}


Expected output(s):

```
$ srun -A PMIU0184 -n 3 ./exercisel9 q1 1 -2 3
srun: job 7294937 queued and waiting for resources
srun: job 7294937 has been allocated resources
mean: 0.666667, variance: 4.22222
```

```
$ srun -A PMIU0110 -n 4 ./exercisel9 q1 1 2 3 4
srun: job 7301859 queued and waiting for resources
srun: job 7301859 has been allocated resources
mean: 2.5, variance: 1.25
```

Part #2.3: Aggregate

In the supplied exercisel9.cpp starter code, implement the aggregate method. The getMeanAndVar method is called on **n** processes of an MPI program. Each process will have a different value (as shown in the adjacent). This method must be implemented to collect/aggregate the values (in order of rank) from all processes into a vector and return the vector in all processes.



Expected output(s):

```
$ srun -A PMIU0110 -n 4 ./exercisel9 q2 1 2 3 4
srun: job 7302122 queued and waiting for resources
srun: job 7302122 has been allocated resources
List[0] = 1
List[1] = 2
List[2] = 3
List[3] = 4
```

```
$ srun -A PMIU0110 -n 4 ./exercisel9 q2 1 5 9 3 7
srun: job 7302677 queued and waiting for resources
srun: job 7302677 has been allocated resources
List[0] = 1
List[1] = 5
List[2] = 9
List[3] = 3
List[4] = 7
```

Part #2.2: Print in-order

In the supplied exercisel9.cpp starter code, implement the printInOrder method. This method is called on all processes of an MPI program. Each process will have a different value supplied as the parameter. Implement this method to print the value (on each process) in the

order of the rank of the process. Hint: barrier in a loop and the i^{th} process prints in the i^{th} iteration of the loop.

Expected output(s):

```
$ srun -A PMIU0110 -n 5 ./exercise19 q3 1 5 9 3 7
srun: job 7302925 queued and waiting for resources
srun: job 7302925 has been allocated resources
Value at rank #0 = "data on 0 is 1_0"
Value at rank #1 = "data on 1 is 1_1"
Value at rank #2 = "data on 2 is 1_2"
Value at rank #3 = "data on 3 is 1_3"
Value at rank #4 = "data on 4 is 1_4"
```

```
$ srun -A PMIU0110 -n 3 ./exercise19 q3 test
srun: job 7302964 queued and waiting for resources
srun: job 7302964 has been allocated resources
Value at rank #0 = "data on 0 is test_0"
Value at rank #1 = "data on 1 is test_1"
Value at rank #2 = "data on 2 is test_2"
```

Part #3: Upload solution to Canvas

Once you have successfully completed and tested the program, submit the following via the Canvas CODE plugin

1. This MS-Word document saved as a PDF file with the convention MUID_Exercise19.pdf.
2. The program you completed as part of this exercise with the source file named with the convention MUID_exercise19.cpp.

Ensure you actually complete the submission process in Canvas (after you accept submission in CODE).